

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 import matplotlib.pyplot as plt,seaborn as sns
```

```
In [2]: 1 traind=pd.read_csv(r"C:\Users\teppa\Downloads\Mobile_Price_Classification_train.csv")
        2 traind
```

```
Out[2]:
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	n_cores	...	px_height	px_width	ram	sc_h
0	842	0	2.2	0	1	0	7	0.6	188	2	...	20	756	2549	9
1	1021	1	0.5	1	0	1	53	0.7	136	3	...	905	1988	2631	17
2	563	1	0.5	1	2	1	41	0.9	145	5	...	1263	1716	2603	11
3	615	1	2.5	0	0	0	10	0.8	131	6	...	1216	1786	2769	16
4	1821	1	1.2	0	13	1	44	0.6	141	2	...	1208	1212	1411	8
...
1995	794	1	0.5	1	0	1	2	0.8	106	6	...	1222	1890	668	13
1996	1965	1	2.6	1	0	0	39	0.2	187	4	...	915	1965	2032	11
1997	1911	0	0.9	1	1	1	36	0.7	108	8	...	868	1632	3057	9
1998	1512	0	0.9	0	4	1	46	0.1	145	5	...	336	670	869	18
1999	510	1	2.0	1	5	1	45	0.9	168	6	...	483	754	3919	19

2000 rows × 21 columns



```
In [3]: 1 testd=pd.read_csv(r"C:\Users\teppa\Downloads\Mobile_Price_Classification_test.csv")
        2 testd
```

```
Out[3]:
```

	id	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	m_dep	mobile_wt	...	pc	px_height	px_width	ram	sc_h
0	1	1043	1	1.8	1	14	0	5	0.1	193	...	16	226	1412	3476	12
1	2	841	1	0.5	1	4	1	61	0.8	191	...	12	746	857	3895	6
2	3	1807	1	2.8	0	1	0	27	0.9	186	...	4	1270	1366	2396	17
3	4	1546	0	0.5	1	18	1	25	0.5	96	...	20	295	1752	3893	10
4	5	1434	0	1.4	0	11	1	49	0.5	108	...	18	749	810	1773	15
...
995	996	1700	1	1.9	0	0	1	54	0.5	170	...	17	644	913	2121	14
996	997	609	0	1.8	1	0	0	13	0.9	186	...	2	1152	1632	1933	8
997	998	1185	0	1.4	0	1	1	8	0.5	80	...	12	477	825	1223	5
998	999	1533	1	0.5	1	0	0	50	0.4	171	...	12	38	832	2509	15
999	1000	1270	1	0.5	0	4	1	35	0.1	140	...	19	457	608	2828	9

1000 rows × 21 columns



```
In [4]: 1 traind.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   battery_power    2000 non-null   int64
1   blue             2000 non-null   int64
2   clock_speed      2000 non-null   float64
3   dual_sim         2000 non-null   int64
4   fc               2000 non-null   int64
5   four_g           2000 non-null   int64
6   int_memory       2000 non-null   int64
7   m_dep            2000 non-null   float64
8   mobile_wt        2000 non-null   int64
9   n_cores          2000 non-null   int64
10  pc               2000 non-null   int64
11  px_height        2000 non-null   int64
12  px_width         2000 non-null   int64
13  ram              2000 non-null   int64
14  sc_h             2000 non-null   int64
15  sc_w             2000 non-null   int64
16  talk_time        2000 non-null   int64
17  three_g          2000 non-null   int64
18  touch_screen     2000 non-null   int64
19  wifi             2000 non-null   int64
20  price_range      2000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 328.2 KB
```

```
In [5]: 1 testd.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 21 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   id              1000 non-null   int64
1   battery_power    1000 non-null   int64
2   blue             1000 non-null   int64
3   clock_speed      1000 non-null   float64
4   dual_sim         1000 non-null   int64
5   fc               1000 non-null   int64
6   four_g           1000 non-null   int64
7   int_memory       1000 non-null   int64
8   m_dep            1000 non-null   float64
9   mobile_wt        1000 non-null   int64
10  n_cores          1000 non-null   int64
11  pc               1000 non-null   int64
12  px_height        1000 non-null   int64
13  px_width         1000 non-null   int64
14  ram              1000 non-null   int64
15  sc_h             1000 non-null   int64
16  sc_w             1000 non-null   int64
17  talk_time        1000 non-null   int64
18  three_g          1000 non-null   int64
19  touch_screen     1000 non-null   int64
20  wifi             1000 non-null   int64
dtypes: float64(2), int64(19)
memory usage: 164.2 KB
```

```
In [28]: 1 x=traind.drop('wifi',axis=1)
2 y=traind["wifi"]
```

```
In [7]: 1 traind['dual_sim'].value_counts()
```

```
Out[7]: dual_sim
1    1019
0     981
Name: count, dtype: int64
```

```
In [8]: 1 m={"three_g":{"yes":1,"No":0}}
2 td=trainind.replace(m)
3 print(td)
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory
0	842	0	2.2	0	1	0	7
1	1021	1	0.5	1	0	1	53
2	563	1	0.5	1	2	1	41
3	615	1	2.5	0	0	0	10
4	1821	1	1.2	0	13	1	44
...
1995	794	1	0.5	1	0	1	2
1996	1965	1	2.6	1	0	0	39
1997	1911	0	0.9	1	1	1	36
1998	1512	0	0.9	0	4	1	46
1999	510	1	2.0	1	5	1	45

	m_dep	mobile_wt	n_cores	...	px_height	px_width	ram	sc_h	sc_w
0	0.6	188	2	...	20	756	2549	9	7
1	0.7	136	3	...	905	1988	2631	17	3
2	0.9	145	5	...	1263	1716	2603	11	2
3	0.8	131	6	...	1216	1786	2769	16	8
4	0.6	141	2	...	1208	1212	1411	8	2
...
1995	0.8	106	6	...	1222	1890	668	13	4
1996	0.2	187	4	...	915	1965	2032	11	10
1997	0.7	108	8	...	868	1632	3057	9	1
1998	0.1	145	5	...	336	670	869	18	10
1999	0.9	168	6	...	483	754	3919	19	4

	talk_time	three_g	touch_screen	wifi	price_range
0	19	0	0	1	1
1	7	1	1	0	2
2	9	1	1	0	2
3	11	1	0	0	2
4	15	1	1	0	1
...
1995	19	1	1	0	0
1996	16	1	1	1	2
1997	5	1	1	0	3
1998	19	1	1	1	0
1999	2	1	1	1	3

[2000 rows x 21 columns]

```
In [9]: 1 x=trainind.drop('wifi',axis=1)
2 y=trainind['wifi']
```

```
In [10]: 1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=42)
3 x_train.shape,x_test.shape
```

Out[10]: ((1600, 20), (400, 20))

```
In [11]: 1 from sklearn.ensemble import RandomForestClassifier
2 rfc=RandomForestClassifier()
3 rfc.fit(x_train,y_train)
```

Out[11]:

RandomForestClassifier

RandomForestClassifier()

```
In [16]: 1 rf=RandomForestClassifier()
```

```
In [36]: 1 params={'max_depth':[2,35,10,20], 'min_samples_leaf':[5,10,20,50,100,200], 'n_estimators':[10,25,30,50,100,200]}
```

```
In [14]: 1 from sklearn.model_selection import GridSearchCV
2 grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring="accuracy")
3 grid_search.fit(x_train,y_train)
```

```
Out[14]: > GridSearchCV
> estimator: RandomForestClassifier
  > RandomForestClassifier
```

```
In [15]: 1 grid_search.best_score_
```

```
Out[15]: 0.5387500000000001
```

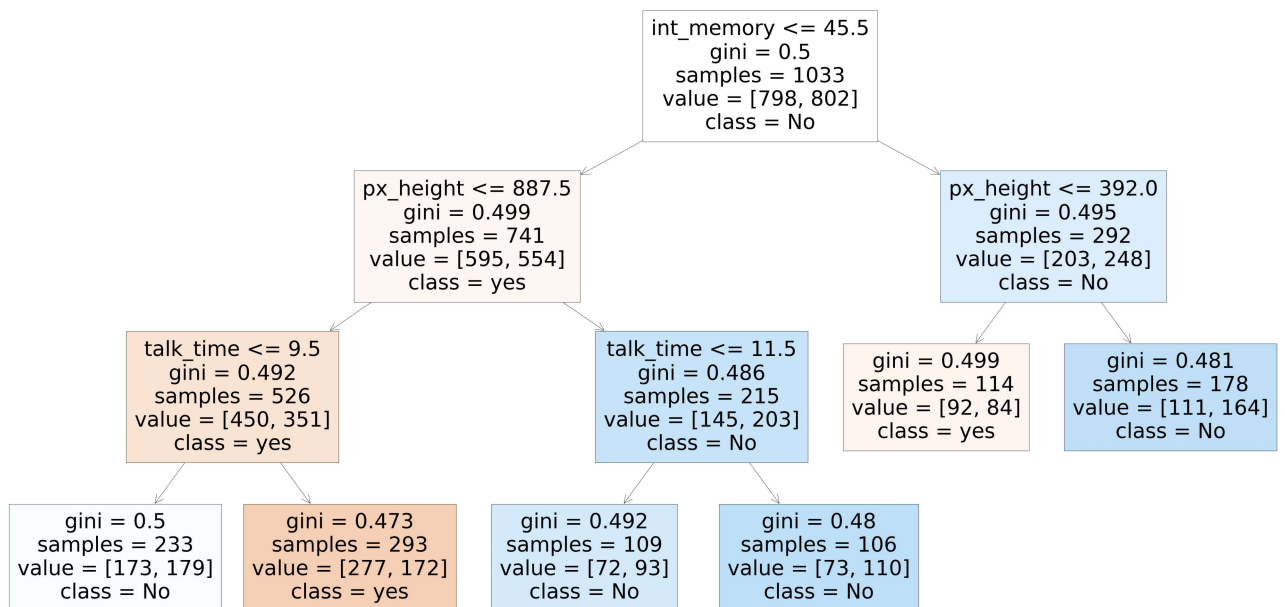
```
In [18]: 1 rf_best=grid_search.best_estimator_
2 print(rf_best)
```

```
RandomForestClassifier(max_depth=3, min_samples_leaf=100, n_estimators=50)
```

```
In [19]: 1 rf_best=grid_search.best_estimator_
2 print(rf_best)
```

```
RandomForestClassifier(max_depth=3, min_samples_leaf=100, n_estimators=50)
```

```
In [21]: 1 from sklearn.tree import plot_tree
2 plt.figure(figsize=(80,40))
3 plot_tree(rf_best.estimators_[5],feature_names=x.columns,class_names=["yes", "No"],filled=True);
```



```
In [23]: 1 rf_best.feature_importances_
```

```
Out[23]: array([0.06249034, 0.0075286 , 0.04696915, 0.00864533, 0.06713393,
0.01704392, 0.06086025, 0.04680869, 0.05561469, 0.01160255,
0.02707365, 0.2379613 , 0.09129798, 0.07013714, 0.04267532,
0.01711047, 0.08518967, 0.00324524, 0.02976883, 0.01084295])
```

```
In [27]: 1 imp_df=pd.DataFrame({"Variance":x_train.columns,"Imp":rf_best.feature_importances_})
2 imp_df.sort_values(by="Imp",ascending=False)
```

Out[27]:

	Variance	Imp
11	px_height	0.237961
12	px_width	0.091298
16	talk_time	0.085190
13	ram	0.070137
4	fc	0.067134
0	battery_power	0.062490
6	int_memory	0.060860
8	mobile_wt	0.055615
2	clock_speed	0.046969
7	m_dep	0.046809
14	sc_h	0.042675
18	touch_screen	0.029769
10	pc	0.027074
15	sc_w	0.017110
5	four_g	0.017044
9	n_cores	0.011603
19	price_range	0.010843
3	dual_sim	0.008645
1	blue	0.007529
17	three_g	0.003245

Test Data

```
In [29]: 1 x=testd.drop('wifi',axis=1)
2 y=testd["wifi"]
```

```
In [30]: 1 testd['dual_sim'].value_counts()
```

```
Out[30]: dual_sim
1      517
0      483
Name: count, dtype: int64
```

```
In [31]: 1 m={"three_g":{"yes":1,"No":0}}
2         td=trainind.replace(m)
3         print(td)
```

	battery_power	blue	clock_speed	dual_sim	fc	four_g	int_memory	
0	842	0	2.2	0	1	0	7	\
1	1021	1	0.5	1	0	1	53	
2	563	1	0.5	1	2	1	41	
3	615	1	2.5	0	0	0	10	
4	1821	1	1.2	0	13	1	44	
...	
1995	794	1	0.5	1	0	1	2	
1996	1965	1	2.6	1	0	0	39	
1997	1911	0	0.9	1	1	1	36	
1998	1512	0	0.9	0	4	1	46	
1999	510	1	2.0	1	5	1	45	

	m_dep	mobile_wt	n_cores	...	px_height	px_width	ram	sc_h	sc_w	
0	0.6	188	2	...	20	756	2549	9	7	\
1	0.7	136	3	...	905	1988	2631	17	3	
2	0.9	145	5	...	1263	1716	2603	11	2	
3	0.8	131	6	...	1216	1786	2769	16	8	
4	0.6	141	2	...	1208	1212	1411	8	2	
...	
1995	0.8	106	6	...	1222	1890	668	13	4	
1996	0.2	187	4	...	915	1965	2032	11	10	
1997	0.7	108	8	...	868	1632	3057	9	1	
1998	0.1	145	5	...	336	670	869	18	10	
1999	0.9	168	6	...	483	754	3919	19	4	

	talk_time	three_g	touch_screen	wifi	price_range
0	19	0	0	1	1
1	7	1	1	0	2
2	9	1	1	0	2
3	11	1	0	0	2
4	15	1	1	0	1
...
1995	19	1	1	0	0
1996	16	1	1	1	2
1997	5	1	1	0	3
1998	19	1	1	1	0
1999	2	1	1	1	3

[2000 rows x 21 columns]

```
In [32]: 1 x=testd.drop('wifi',axis=1)
2         y=testd['wifi']
```

```
In [33]: 1 from sklearn.model_selection import train_test_split
2         x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2,random_state=42)
3         x_train.shape,x_test.shape
```

Out[33]: ((800, 20), (200, 20))

```
In [34]: 1 from sklearn.ensemble import RandomForestClassifier
2         rfc=RandomForestClassifier()
3         rfc.fit(x_train,y_train)
```

Out[34]:

RandomForestClassifier

RandomForestClassifier()

```
In [35]: 1 rf=RandomForestClassifier()
```

```
In [37]: 1 params={'max_depth':[2,35,10,20], 'min_samples_leaf':[5,10,20,50,100,200], 'n_estimators':[10,25,30,50,100,200]}
```

```
In [38]: 1 from sklearn.model_selection import GridSearchCV
2 grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring="accuracy")
3 grid_search.fit(x_train,y_train)
```

```
Out[38]: > GridSearchCV
> estimator: RandomForestClassifier
  > RandomForestClassifier
```

```
In [39]: 1 grid_search.best_score_
```

```
Out[39]: 0.53
```

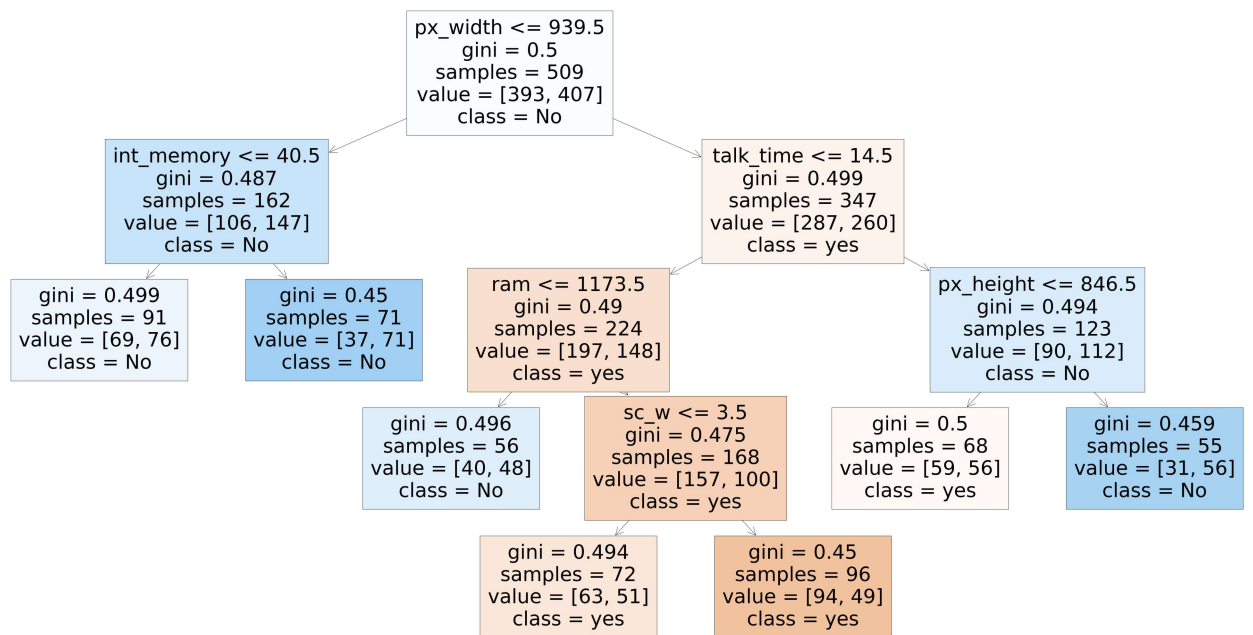
```
In [40]: 1 rf_best=grid_search.best_estimator_
2 print(rf_best)
```

```
RandomForestClassifier(max_depth=35, min_samples_leaf=50, n_estimators=50)
```

```
In [41]: 1 rf_best=grid_search.best_estimator_
2 print(rf_best)
```

```
RandomForestClassifier(max_depth=35, min_samples_leaf=50, n_estimators=50)
```

```
In [42]: 1 from sklearn.tree import plot_tree
2 plt.figure(figsize=(80,40))
3 plot_tree(rf_best.estimators_[5],feature_names=x.columns,class_names=["yes", "No"],filled=True);
```



```
In [43]: 1 rf_best.feature_importances_
```

```
Out[43]: array([0.04405736, 0.08051273, 0.01745645, 0.05571116, 0.01453586,
0.06371533, 0.03197881, 0.06969091, 0.01912393, 0.11489556,
0.01938361, 0.04267464, 0.06589819, 0.09438956, 0.08943337,
0.02961427, 0.02508144, 0.10637251, 0.00143634, 0.01403797])
```

```
In [44]: 1 imp_df=pd.DataFrame({"Variance":x_train.columns,"Imp":rf_best.feature_importances_})
2 imp_df.sort_values(by="Imp",ascending=False)
```

Out[44]:

	Variance	Imp
9	mobile_wt	0.114896
17	talk_time	0.106373
13	px_width	0.094390
14	ram	0.089433
1	battery_power	0.080513
7	int_memory	0.069691
12	px_height	0.065898
5	fc	0.063715
3	clock_speed	0.055711
0	id	0.044057
11	pc	0.042675
6	four_g	0.031979
15	sc_h	0.029614
16	sc_w	0.025081
10	n_cores	0.019384
8	m_dep	0.019124
2	blue	0.017456
4	dual_sim	0.014536
19	touch_screen	0.014038
18	three_g	0.001436

```
In [ ]: 1
```